

ex-actor-critic

November 26, 2024

1 Actor-Critic Algorithm

In this notebook, you'll code The Actor-Critic Algorithm from scratch: .

Actor-Critic algorithm is a *Policy-based method* that aims to reduce the variance of the Reinforce algorithm and train our agent faster and better by using a combination of Policy-Based and Value-Based methods

To test its robustness, we're going to train it in Cartpole-v1 environment

1.1 Import the packages

```
[ ]: !pip install swig
      !pip install gymnasium[box2d]
      !pip install gymnasium
```

```
[ ]: import numpy as np
      import pandas as pd
      from collections import deque
      import matplotlib.pyplot as plt

      # PyTorch
      import torch
      import torch.nn as nn
      import torch.nn.functional as F
      import torch.optim as optim
      from torch.distributions import Categorical

      # Gym
      import gymnasium as gym
```

1.2 Check if we have a GPU

- Let's check if we have a GPU device:cuda0

```
[ ]: device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
```

2 Agent: Playing CartPole-v1

2.0.1 The CartPole-v1 environment

A pole is attached by an un-actuated joint to a cart, which moves along a frictionless track. The pendulum is placed upright on the cart and the goal is to balance the pole by applying forces in the left and right direction on the cart.

So, we start with CartPole-v1. The goal is to push the cart left or right **so that the pole stays in the equilibrium.**

The episode ends if: - The pole Angle is greater than $\pm 12^\circ$ - Cart Position is greater than ± 2.4 - Episode length is greater than 500

We get a reward of +1 every timestep the Pole stays in the equilibrium.

```
[ ]: env_id = "CartPole-v1"
      #env_id="LunarLander-v2"
      # Create the env
      env = gym.make(env_id)

      # Get the state space and action space
      s_size = env.observation_space.shape[0]
      a_size = env.action_space.n
```

2.1 Let's build the A2C algo

```
[ ]: #Using a neural network to learn our actor (policy) parameters

class Actor(nn.Module):
    def __init__(self, s_size, a_size, h_size):
        super(Actor, self).__init__()
        self.fc1 = nn.Linear(s_size, h_size)
        self.fc2 = nn.Linear(h_size, a_size)

    def forward(self, x):
        x = F.relu(self.fc1(x))
        x = self.fc2(x)
        return F.softmax(x, dim=1)

    def act(self, state):

        probs = self.forward(state)
        m = Categorical(probs)
        action = m.sample()
        return action.item(), m.log_prob(action)
```

```
[ ]: #Using a neural network to learn state value
class Critic(nn.Module):

    #Takes in state
    def __init__(self, s_size, h_size):
        super(Critic, self).__init__()

        # two fully connected layers
        # add code here
        # add code here

    def forward(self, x):

        #input layer
        x = self.input_layer(x)

        #activation relu
        x = F.relu(x)

        #get state value
        state_value = self.output_layer(x)

        return state_value
```

3 Building the parts of our algorithm

The main steps for building a A2C Algo are: 1. Generates a trajectory 2. Compute the discounted returns 3. Standardization of the returns 4. Train critic network 5. Train actor network

```
[ ]: def generate_trajectory(actor, critic, max_t):
    saved_log_probs = []
    rewards = []
    state_values=[]

    state = env.reset()
    for t in range(max_t):
        state=torch.from_numpy(state).float().unsqueeze(0).to(device)

        action, log_prob = # add code here

        # get the state value from th critic network
        state_val= # add code here

        next_state, reward, done, _ = env.step(action)

        # add te obtained results to their relative lists ==>
        ↪ saved_log_probs, rewards, state_values
```

```

        # add code here
        # add code here
        # add code here

    state=next_state

    if done:

        break

    return saved_log_probs, rewards, state_values

```

```

[ ]: def computer_cumulative_reward(rewards, max_t,gamma):

    returns = deque(maxlen=max_t)
    n_steps = len(rewards)
    for t in range(n_steps)[::-1]:
        disc_return_t = (returns[0] if len(returns)>0 else 0)
        returns.appendleft( rewards[t]+gamma*disc_return_t)
    return returns

```

```

[ ]: def returns_standardization(returns):
    eps = np.finfo(np.float32).eps.item()
    ## eps is the smallest representable float, which is
    # added to the standard deviation of the returns to avoid numerical
    ↪ instabilities
    returns = torch.tensor(returns).to(device)
    returns = (returns - returns.mean()) / (returns.std() + eps)

    return returns

```

```

[ ]: def train_actor(actorOptimizer,saved_log_probs, returns,state_values):

    state_values= torch.stack(state_values).squeeze()

    #calculate Advantage for actor
    advantages = # add code here

    #convect the advantages to a tensor
    advantages = # add code here

    actor_loss = []
    # compute the actor loss
    # add code here

```

```

actor_loss = torch.cat(actor_loss).sum()
# Backpropagate actor
actorOptimizer.zero_grad()
actor_loss.backward()
actorOptimizer.step()

```

```

[ ]: def train_critic(criticOptimizer, returns, state_values):
    state_values= torch.stack(state_values).squeeze()

    critic_loss=#add code here

    # Backpropagate critic
    criticOptimizer.zero_grad()
    critic_loss.backward()
    criticOptimizer.step()

```

3.1 Merge all functions into the Actor_Critic method

You will now see how the overall **A2C Algo** is structured by putting together all the building blocks (functions implemented in the previous parts) together, in the right order.

```

[ ]: def Actor_Critic(actor, critic, actorOptimizer, criticOptimizer,
    ↪ n_training_episodes, max_t, gamma, print_every):
    # Help us to calculate the score during the training
    scores_deque = deque(maxlen=100)
    scores = []

    for i_episode in range(1, n_training_episodes+1):

        # Generate an episode
        #add code here
        scores_deque.append(sum(rewards))
        scores.append(sum(rewards))

        # calculate the return
        returns= computer_cumulative_reward(rewards, max_t, gamma)

        ## standardization of the returns is employed to make training more
        ↪ stable
        returns=returns_standardization(returns)

        # Train the Critic network
        #add code here

        # Train the Actor network
        #add code here

```

```

        if i_episode % print_every == 0:
            print('Episode {} \t Average Score: {:.2f}'.format(i_episode, np.
↪ mean(scores_deque)))

    return scores

```

3.2 Train it

- We're now ready to train our agent.
- But first, we define a variable containing all the training hyperparameters.
- You can change the training parameters (and should)

```

[ ]: cartpole_hyperparameters = {
    "h_size": 64,
    "n_training_episodes": 1000,
    "n_evaluation_episodes": 10,
    "max_t": 1000,
    "gamma": 1.0,
    "lr": 1e-2,
    "env_id": env_id,
    "state_space": s_size,
    "action_space": a_size,
}

```

```

[ ]: # Create actor and place it to the device
    cartpole_actor = #add code here

    cartpole_actorOptimizer = #add code here

```

```

[ ]: # Create critic and place it to the device
    cartpole_critic = #add code here
    cartpole_criticOptimizer = #add code here

```

```

[ ]: scores = Actor_Critic(cartpole_actor,
    ↪
    ↪ cartpole_critic, cartpole_actorOptimizer, cartpole_criticOptimizer,
        cartpole_hyperparameters["n_training_episodes"],
        cartpole_hyperparameters["max_t"],
        cartpole_hyperparameters["gamma"],
        100)

```

```

[ ]: scores = pd.Series(scores, name="scores_Actor")
    scores.describe()

```

```
[ ]: fig, ax = plt.subplots(1, 1)
_ = scores.plot(ax=ax, label="scores_Actor")
_ = (scores.rolling(window=100)
      .mean()
      .rename("Rolling Average")
      .plot(ax=ax))
ax.legend()
_ = ax.set_xlabel("Episode Number")
_ = ax.set_ylabel("scores_Actor")
```

```
[ ]:
```